

Mini Project

Contents

- [Exploratory Data Analysis](#)
- [Data preparation](#)
- [Model Helper Functions](#)
- [XGBoost Classifier](#)
 - [Feature Importance](#)
- [Multilayer Perceptron](#)
- [Accuracy vs Number of Features](#)
- [PCA](#)

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import xgboost
from xgboost import XGBClassifier
from sklearn.model_selection import train_test_split
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.neural_network import MLPClassifier
import shap
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.decomposition import PCA
from sklearn import preprocessing
from matplotlib import cm
```

```
In [2]: data = pd.read_csv("../Data/DataSet_HAR.csv")
subjects = data['subject'].drop_duplicates()
data.set_index('subject', inplace = True)
```

EDA

```
In [3]: data.shape
```

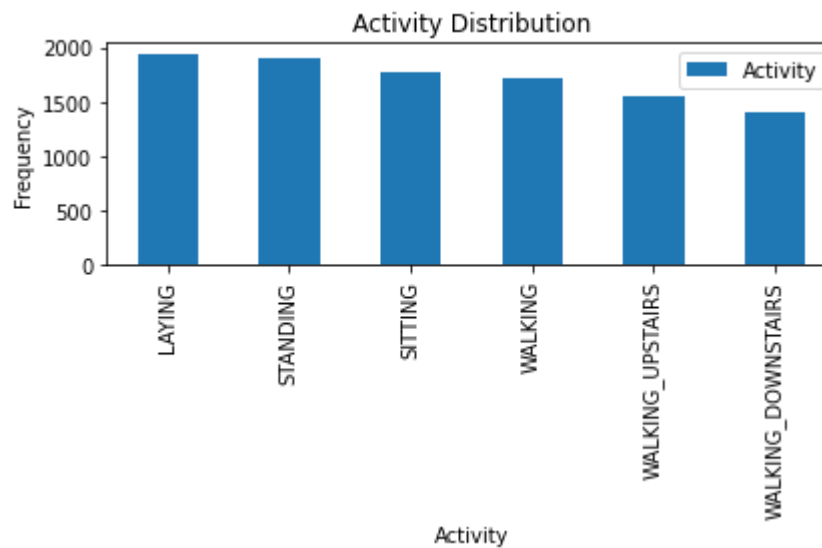
```
Out[3]: (10299, 562)
```

```
In [4]: data.isnull().sum().sum()
```

```
Out[4]: 0
```

```
In [5]: fig = plt.figure(figsize=(15,8))
data[['Activity']].apply(pd.value_counts).plot(kind='bar')
plt.xlabel('Activity')
plt.ylabel('Frequency')
plt.title('Activity Distribution')
#plt.show()
plt.tight_layout()
plt.savefig("Activity Distribution")
```

<Figure size 1080x576 with 0 Axes>



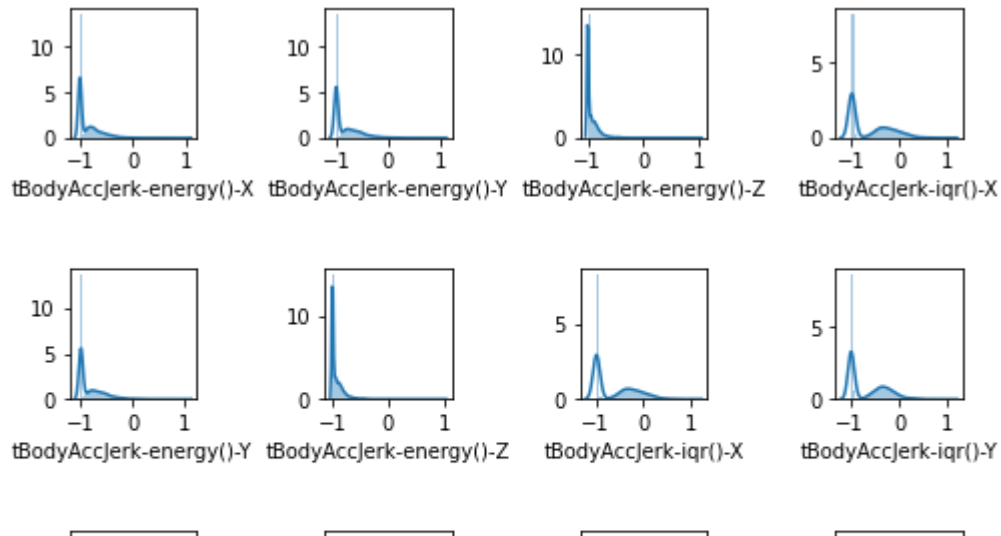
```
In [6]: columns = data.columns
columns = columns[0:561]
```

In [7]:

```

for i in range(47):
    fig, axs = plt.subplots(3,4)
    fig.subplots_adjust(hspace=1,wspace = 1)
    fig.set_size_inches(8,6)
    for j in range(3):
        for k in range(4):
            sns.distplot(data[columns[i*12 + j + k]], ax = axs[j,k])
    #fig.savefig('../Plots/dist_' + str(i) + '.png', dpi=100)

```



In [8]:

```
for i in range(len(columns[0:47])):
    fig, axs = plt.subplots(3,4)
    fig.subplots_adjust(hspace=1,wspace = 1)
    fig.set_size_inches(8,6)
    for j in range(3):
        for k in range(4):
            ax = axs[j,k]
            data.groupby('Activity')[columns[i]].hist(alpha=0.4, ax=ax)
            ax.set_title(columns[i] + '_hist')
            ax.legend()

# fig.savefig('../Plots/classwise_dist_' + str(i) + '.png', dpi=100)
```

No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.

No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.
No handles with labels found to put in legend.

In [9]:

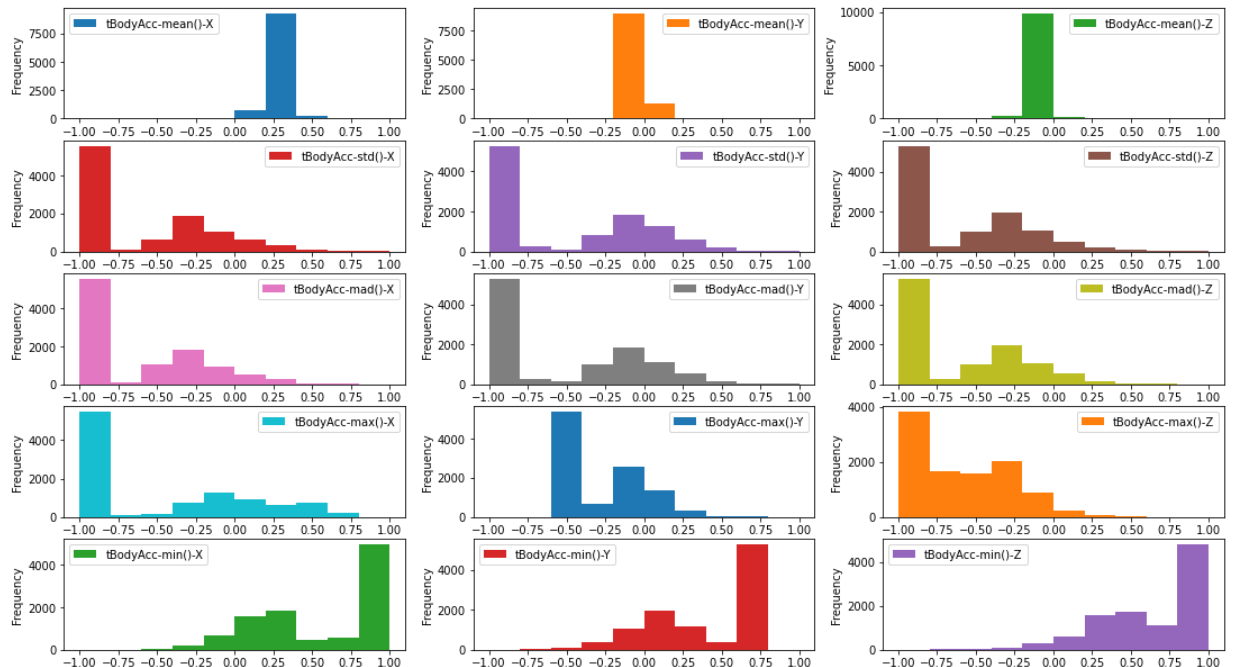
```
fig, axes = plt.subplots(nrows=5, ncols=3)
data[data.columns[0:15]].plot(subplots = True, ax = axes, kind = 'hist', bins = 10)
fig.set_size_inches(18.5, 10.5)
plt.savefig('body_Acc.png')
```

The rowNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use `ax.get_subplotspec().rowspan.start` instead.

The colNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use `ax.get_subplotspec().colspan.start` instead.

The rowNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use `ax.get_subplotspec().rowspan.start` instead.

The colNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use `ax.get_subplotspec().colspan.start` instead.



In [10]:

```
def save_plot(column):
    fig,ax = plt.subplots()
    g = data.groupby('Activity')


    num_groups = g.ngroups

    for i, group in g:
        group[column].hist(alpha=0.7, ax =ax , label=i)


    ax.legend()
    plt.xlabel(column)
    plt.ylabel("Frequency")
    plt.title(column + " Distribution")
    plt.savefig(column + ".png")
```

In [11]:

```
l = ['tGravityAcc-min()-X', 'tBodyAcc-max()-X',
     'fBodyAcc-bandsEnergy()-1,8', 'tBodyGyroJerk-iqr()-X',
     'tGravityAcc-max()-Y', 'tBodyGyro-max()-X', 'tBodyAccJerk-max()-X',
     'tGravityAcc-mean()-X', 'tBodyAcc-correlation()-X,Y',
     'fBodyGyro-bandsEnergy()-9,16']
for i in l:
    save_plot(i)
```



```
In [12]: save_plot(data.columns[4])
```



Data Prep For model

```
In [13]: def data_prep(train_set, test_set, columns):  
  
    train_set['Activity'] = pd.factorize(train_set['Activity'], sort = True)[0]  
    test_set['Activity'] = pd.factorize(test_set['Activity'], sort = True)[0]  
    X_train = train_set[columns]  
    X_test = test_set[columns]  
    y_train = train_set['Activity']  
    y_test = test_set['Activity']  
  
    return (X_train, X_test, y_train, y_test)
```

```
In [14]: def train_test(X, Y):  
    test_size = 0.3  
    seed = 123  
    X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=test_size,  
                                                         random_state=seed)  
    return (X_train, X_test, y_train, y_test)
```

```
In [15]: data_org = data.copy()
```

```
In [16]: X = data[list(data.columns)[0:561]]
```

```
In [17]: Y = data[list(data.columns)[561]]

encoder = preprocessing.LabelEncoder()

# encoding train labels
encoder.fit(Y)
```

```
Out[17]: LabelEncoder()
```

```
In [18]: X_train, X_test, y_train, y_test = train_test(X, encoder.transform(Y))
```

```
In [19]: res = {}
for cl in encoder.classes_:
    res.update({cl:encoder.transform([cl])[0]})
res
```

```
Out[19]: {'LAYING': 0,
'SITTING': 1,
'STANDING': 2,
'WALKING': 3,
'WALKING_DOWNSTAIRS': 4,
'WALKING_UPSTAIRS': 5}
```

Model Helper Functions

```
In [20]: def train_model(X, Y, model):

    model.fit(X, Y)

    return model
```

```
In [21]: def get_confusion_matrix(y_pred, y_true):

    accuracy = accuracy_score(y_true, y_pred)
    cm = confusion_matrix(y_true, y_pred, labels = [0, 1, 2, 3, 4, 5])

    return (accuracy, cm)
```

XGboost Classifier

```
In [22]: xgb_model = XGBClassifier()
```



```
In [23]: model = train_model(X_train, y_train, xgb_model)
y_test_pred = model.predict(X_test)
y_train_pred = model.predict(X_train)
```

```
In [24]: acc_train, cm = get_confusion_matrix(y_train_pred, y_train)
acc_test, cm = get_confusion_matrix(y_test_pred, y_test)
```

```
In [25]: print("Training Accuracy", acc_train)
print("Testing Accuracy", acc_test)
```

```
Training Accuracy 0.999445138021917
Testing Accuracy 0.9828478964401295
```

```
In [26]: print(cm)
```

```
[[543  0  0  0  0  0]
 [ 0 564 23  0  0  0]
 [ 0 14 544  0  0  0]
 [ 0  0  0 503  0  3]
 [ 0  0  0  1 438  9]
 [ 0  0  0  2  1 445]]
```

```
In [27]: print(classification_report(list(encoder.inverse_transform(y_test_pred)), list(enc
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	543
SITTING	0.96	0.98	0.97	578
STANDING	0.97	0.96	0.97	567
WALKING	0.99	0.99	0.99	506
WALKING_DOWNSTAIRS	0.98	1.00	0.99	439
WALKING_UPSTAIRS	0.99	0.97	0.98	457
accuracy			0.98	3090
macro avg	0.98	0.98	0.98	3090
weighted avg	0.98	0.98	0.98	3090

Feature importance

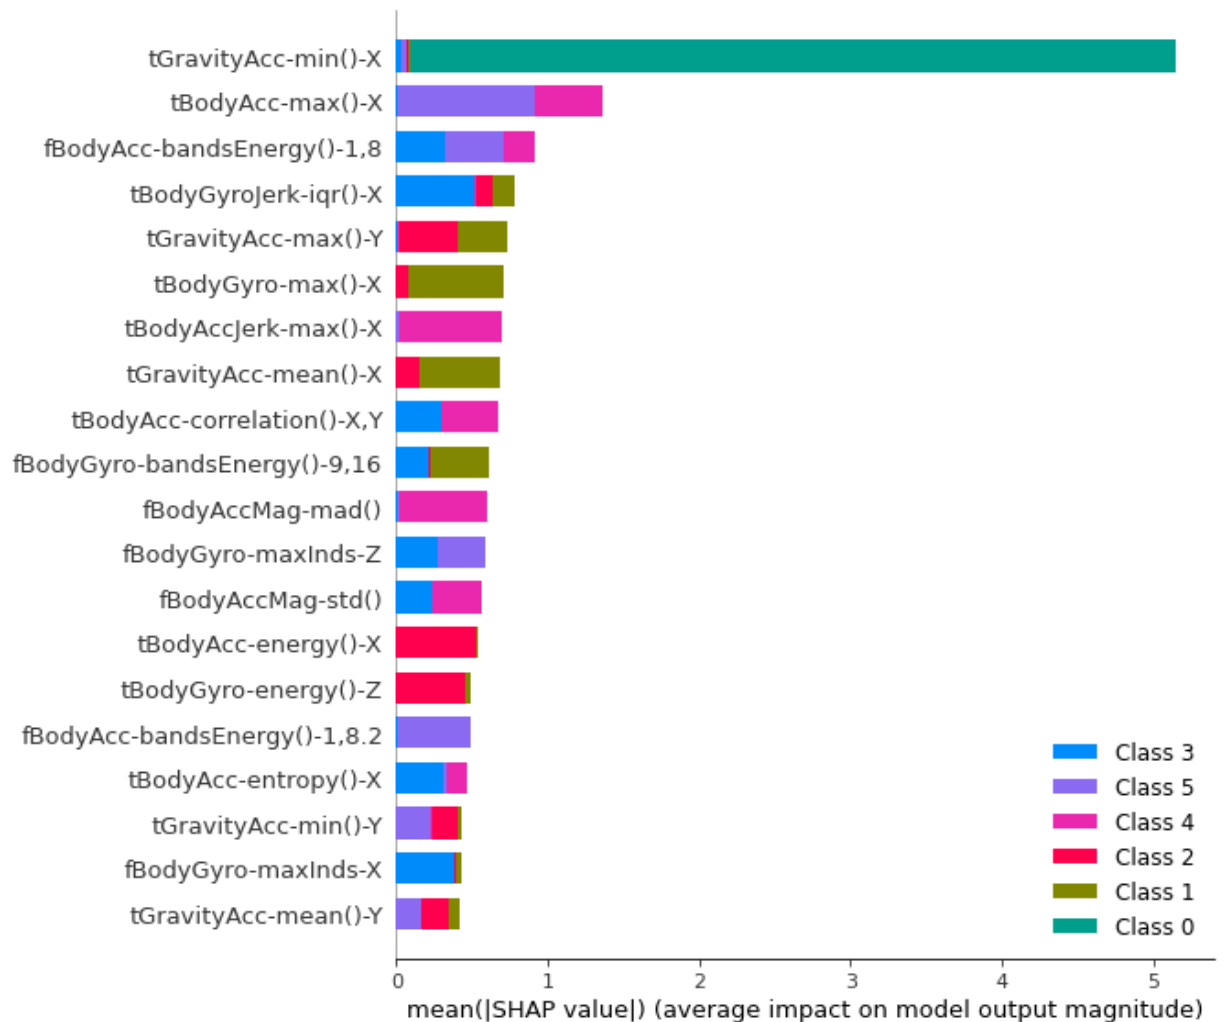
```
In [28]: shap.initjs()
```



```
In [29]: explainer = shap.TreeExplainer(model)
shap_values = explainer.shap_values(X_test)
```

Setting feature_perturbation = "tree_path_dependent" because no background data was given.

```
In [30]: # Plot summary_plot
shap.summary_plot(shap_values, X_test)
plt.savefig("Feature Importance.png")
```



<Figure size 432x288 with 0 Axes>

```
In [31]: vals= np.abs(shap_values).mean(0)
feature_importance = pd.DataFrame(list(zip(X_test.columns, sum(vals))), columns=
feature_importance.sort_values(by=['feature_importance_vals'], ascending=False,i
feature_importance.to_csv('feature importance XgBoost.csv')
```

```
In [32]: feature_importance = feature_importance['col_name'].values
```

```
In [33]: feature_importance[:10]
```

```
Out[33]: array(['tGravityAcc-min()-X', 'tBodyAcc-max()-X',
                'fBodyAcc-bandsEnergy()-1,8', 'tBodyGyroJerk-iqr()-X',
                'tGravityAcc-max()-Y', 'tBodyGyro-max()-X', 'tBodyAccJerk-max()-X',
                'tGravityAcc-mean()-X', 'tBodyAcc-correlation()-X,Y',
                'fBodyGyro-bandsEnergy()-9,16'], dtype=object)
```

MLP Classifier

```
In [34]: mlp_model = MLPClassifier(random_state=1, max_iter=300)
```

```
In [35]: model = train_model(X_train, y_train, mlp_model)
y_test_pred = model.predict(X_test)
y_train_pred = model.predict(X_train)
```

```
In [36]: acc_train, cm = get_confusion_matrix(y_train_pred, y_train)
acc_test, cm = get_confusion_matrix(y_test_pred, y_test)
```

```
In [37]: print("Training Accuracy", acc_train)
print("Testing Accuracy", acc_test)
```

Training Accuracy 0.9933416562630045
Testing Accuracy 0.983495145631068

```
In [38]: print(cm)
```

```
[[543  0  0  0  0  0]
 [ 0 551 36  0  0  0]
 [ 0 15 543  0  0  0]
 [ 0  0  0 506  0  0]
 [ 0  0  0  0 448  0]
 [ 0  0  0  0  0 448]]
```


```
In [39]: print(classification_report(list(encoder.inverse_transform(y_test_pred)), list(en
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	543
SITTING	0.94	0.97	0.96	566
STANDING	0.97	0.94	0.96	579
WALKING	1.00	1.00	1.00	506
WALKING_DOWNSTAIRS	1.00	1.00	1.00	448
WALKING_UPSTAIRS	1.00	1.00	1.00	448
accuracy			0.98	3090
macro avg	0.99	0.99	0.99	3090
weighted avg	0.98	0.98	0.98	3090

Identifying Accuracy vs Number of Columns required

```
In [40]: accuracy = []  
  
for i in range(1,50):  
    columns = feature_importance[0:i]  
    columns = list(columns)  
    xgb_model = XGBClassifier()  
    model = train_model(X_train[columns], y_train, xgb_model)  
    y_pred = model.predict(X_test[columns])  
    acc, cm = get_confusion_matrix(y_pred, y_test)  
    accuracy.append(acc)
```

```
In [41]: plt.plot( list(range(1,25)), accuracy[0:24])  
plt.xlabel("Number of Features")  
plt.ylabel("Accuracy")  
plt.title("Accuracy vs Number of Features")  
#plt.show()  
plt.savefig("Accuracy vs Numbrt of Features")
```



```
In [42]: for i,v in enumerate(accuracy):  
         print(i,v)
```

```
0 0.47896440129449835  
1 0.7692556634304207  
2 0.7828478964401294  
3 0.8543689320388349  
4 0.9100323624595469  
5 0.9187702265372168  
6 0.933009708737864  
7 0.9336569579288025  
8 0.9411003236245955  
9 0.941747572815534  
10 0.9478964401294498  
11 0.9540453074433657  
12 0.9546925566343042  
13 0.9524271844660194  
14 0.9576051779935275  
15 0.9576051779935275  
16 0.959546925566343  
17 0.962135922330097  
18 0.9644012944983819  
19 0.9640776699029127  
20 0.9644012944983819  
21 0.9669902912621359  
22 0.9711974110032362  
23 0.9718446601941747  
24 0.9724919093851133  
25 0.9744336569579288  
26 0.974757281553398  
27 0.9744336569579288  
28 0.9766990291262136  
29 0.976051779935275  
30 0.9766990291262136  
31 0.9783171521035599  
32 0.9770226537216828  
33 0.9776699029126213  
34 0.9783171521035599  
35 0.9786407766990292  
36 0.9773462783171522  
37 0.9770226537216828  
38 0.9773462783171522  
39 0.9792880258899677  
40 0.9799352750809062  
41 0.9796116504854369  
42 0.9799352750809062  
43 0.9802588996763754  
44 0.9818770226537217  
45 0.9802588996763754  
46 0.9809061488673139  
47 0.9796116504854369  
48 0.9783171521035599
```

PCA

```
In [43]: pca = PCA( svd_solver='full')
```

```
In [44]: pca.fit(X_train)
```


```
Out[44]: PCA(copy=True, iterated_power='auto', n_components=None, random_state=None,  
          svd_solver='full', tol=0.0, whiten=False)
```

```
In [45]: pca_sum = pca.explained_variance_ratio_.cumsum()
```

```
In [46]: for i,val in enumerate(pca_sum):  
         print(i,val)
```

```
69 0.9518033138539366  
70 0.9526946136025539  
71 0.9535716205535273  
72 0.9544081457748887  
73 0.9552360403290069  
74 0.9560545172060196  
75 0.9568663939722262  
76 0.9576504098500372  
77 0.9584195689389133  
78 0.9591849146513712  
79 0.9599253780408631  
80 0.9606512129230305  
81 0.9613672119202841  
82 0.9620600125353189  
83 0.9627485283168615  
84 0.9634250107566441  
85 0.9640967379498053  
86 0.9647457067666348  
87 0.9653833265162237  
88 0.9660150270762657
```


```
In [47]: plt.plot(pca_sum)
plt.xlabel("Number of Components")
plt.ylabel("Explained Variance")
plt.title("PCA explained variables vs Number of Features")
#plt.show()
plt.savefig("Explained Variance.png")
```



```
In [48]: accuracy = []

for i in range(1,30):
    xgb_model = XGBClassifier()
    model = train_model(pca.transform(X_train)[:i], y_train, xgb_model)
    y_pred = model.predict(pca.transform(X_test)[:i])
    acc, cm = get_confusion_matrix(y_pred, y_test)
    accuracy.append(acc)
```

```
In [52]: plt.plot(list(range(1,30)), accuracy[0:29])
plt.xlabel("Number of PCA Components")
plt.ylabel("Accuracy")
plt.title("Accuracy vs Number of PCA Components")
#plt.show()
plt.savefig("Accuracy vs Number of PCA Components")
```




```
In [50]: for i,val in enumerate(accuracy):  
         print(i,val)
```

```
0 0.49093851132686084  
1 0.6401294498381876  
2 0.7779935275080906  
3 0.7844660194174757  
4 0.8158576051779936  
5 0.8375404530744337  
6 0.8715210355987055  
7 0.8776699029126214  
8 0.8844660194174757  
9 0.8867313915857605  
10 0.8909385113268609  
11 0.8980582524271845  
12 0.8977346278317152  
13 0.9019417475728155  
14 0.9048543689320389  
15 0.911326860841424  
16 0.9103559870550162  
17 0.9132686084142395  
18 0.9148867313915857  
19 0.9158576051779935  
20 0.9161812297734628  
21 0.9187702265372168  
22 0.917799352750809  
23 0.9194174757281554  
24 0.9213592233009709  
25 0.9216828478964402  
26 0.9210355987055017  
27 0.9210355987055017  
28 0.9197411003236245
```